

bischeck

-

A Business Activity Monitoring Server

Version 0.4.3

2013-03-27

Contents

1	Introduction	5
2	Service configuration	8
3	Services and service items classes	9
3.1	Service classes	9
3.1.1	JDBC	9
3.1.2	Livestatus	9
3.1.3	Bischeck cache	9
3.1.4	Shell	10
3.2	Service item classes	10
3.2.1	SQL	10
3.2.2	Livestatus	10
3.2.3	Cache calculations	10
3.2.4	Check commands	11
4	Date formatting	12
5	Service item cache - LastStatusCache	13
5.1	Methods to retrieve cache data	13
5.1.1	Retrieve single cache data element by index	13
5.1.2	Retrieve single cache data element by time	13
5.1.3	Retrieve a list of cache data elements by index	13
5.1.4	Retrieve a list of cache data elements by time range	14
5.2	Mathematical expressions on cache data	14
5.2.1	JEP standard functions	14
5.2.2	JEP extensions	15
5.3	Null value in the cache	15
6	Threshold configuration	18
6.1	Twenty4HourThreshold class	18
6.1.1	Period definition	18
6.1.2	Calculation definition	19
6.1.3	Hours and threshold definition	19

7	Server integration	21
7.1	Nagios integration with NSCA, NRDP and Livestatus	21
7.1.1	NSCA configuration	22
7.1.2	NRDP configuration	22
7.1.3	Livestatus configuration	23
7.1.4	Nagios state on null	23
7.2	OpenTSDB	23
7.3	Graphite	24
8	Bischeck configuration files	26
8.1	Naming standard	26
8.2	bischeck.xml	26
8.3	properties.xml	28
8.4	urlservice.xml	29
8.5	server.xml	30
8.6	24thresholds.xml	31
9	Service scheduling	35
9.1	Interval scheduling	35
9.2	Cron based scheduling	35
9.3	Service relation scheduling	35
10	Configuration tools	36
11	Bischeck internal surveillance	37
11.1	Bischeck MBeans	37
12	Building Bischeck	39
12.1	Jar customization	39
12.2	Developing with Bischeck	39
13	Installation	40
13.1	Upgrading	40
13.2	Getting started	41
13.3	Integration with pnp4nagios	41
13.4	Logging	41
14	Command line utilities	42
14.1	Run Bischeck	42
14.2	List Bischeck configuration	42
14.3	Validating configuration files	43
14.4	Twenty4HourThreshold testing	43

15 Releases	44
15.1 Release 0.4.3 - 2013-03-27	44
15.1.1 New feature	44
15.1.2 Bugs fixed and important issues	44
15.1.3 Upgrading	44
15.2 Release 0.4.2 - 2012-12-21	45
15.2.1 New feature	45
15.2.2 Bugs fixed and important issues	46
15.2.3 Upgrading	46
15.3 Release 0.4.1 - 2012-10-01	46
15.3.1 New feature	46
15.3.2 Bugs fixed and important issues	47
15.3.3 Upgrading	47
15.4 Release 0.4.0 - 2012-08-31	47
15.4.1 New feature	47
15.4.2 Bugs fixed and important issues	49
15.4.3 Upgrading	50
15.5 Release 0.3.3 - 2011-11-14	50
15.5.1 New features	50
15.5.2 Bugs fixed and important issues	50
15.6 Release 0.3.2 - 2011-07-29	50
15.6.1 New features	50
15.6.2 Bugs fixed and important issues	51
15.7 Release 0.3.1 - 2011-04-08	51
15.7.1 New features	51
15.7.2 Bugs fixed and important issues	52
15.7.3 Upgrade issues	52
15.8 Release 0.3.0 - 2011-03-03	53
15.8.1 New features	53
15.8.2 Bugs fixed and important issues	53
16 System requirements	54
17 Bischeck license	56
18 Bug reports and feature requests	57
19 Credits	58

1 Introduction

Monitoring and surveillance is today common in any IT operation. The market is mature and a number of excellent solutions exist, especially in the open source market, like Nagios. But when it comes to monitoring of business applications, the maturity is not at the same level. Business operations teams seldom have any monitoring solutions to understand what goes on in their business processes and applications. Business intelligence solutions can often tell what happened last month, but from a more tactical operations view there is a gap. So why do not business operations team implement the same surveillance structure or use the existing IT operations infrastructure to understand *“why the number of orders have dropped under the current threshold”*?

What is the difference between business activity monitoring, BAM, and traditional IT infrastructure monitoring? Basically it is the same. Some entity to measure and for that entity to define threshold levels that will generate warnings or critical alarms. The demands we have seen in addition to the basic functionality are the following features provided by Bischeck:

- Support for dynamic and adaptive thresholds.
- Support dynamic configuration of thresholds depending of time of the day and day of week or month.
- Support thresholds based on fixed values as well as thresholds based on measured values from other services. This is typical for process driven businesses.
- Support for a large number of mathematical functions used to calculate thresholds.
- Data cache for all measured service items that can be retrieved based on index or time. Cache data is the main source to create dynamic and adaptive thresholds.
- Support for multiple scheduling schemas per service. This enables a fine grained control of when a service should be run. Bischeck supports interval scheduling, cron scheduling and service relation based scheduling.
- Configuration of services based on multiple measured values of other services, which could be described as "virtual" entities.
- Support for date macros in the execution statements of a service, typically used in a *where* clause when selecting from databases.
- Support for a multitude of ways to connect to source systems by allowing custom developed service connection methods.

- Support for custom threshold classes to enable any way to define threshold logic.
- Integration with multiple monitoring and surveillance tools, like any Nagios like systems that supports the Nagios Service Check Acceptor (NSCA) and/or the Nagios Remote Data Processor (NRDP)[New in 0.4.2] protocol¹, OpenTSDB, Graphite [New in 0.4.2] and more.[New in 0.4.0]

Let us look at 3 examples that can be solved by using Bischeck:

Example 1 – Monitor the number of orders received during the day

The order management application receives orders 24 hours a day during Monday to Friday. The total aggregated number of orders varies as a function of the time of the day. The business expects to have a total number of orders of 1500 at 13:00, at 14:00 the order count should be 2300, at 15:00 it should be 3400, etc. Between every hour the business requirement is to interpolate the order rate according to a linear function between the defined hour rates. This means that the threshold at 13:20 is $(2300-1500)*20/60+1500 = 1767$. The warning alarm level should be between 90% and 70% of the threshold and critical alarm if the measured value is below 70% of the threshold.

Example 2 – Monitor the number of created invoices in relation to the number of received orders

The invoice system should invoice at least 80% of the daily incoming orders in the same day with a maximum of one hour delay. This means that the measured value of orders with one hour delay must be used as a threshold for the number of created invoices.

Example 3 – Monitor the current number of orders and if the incoming flow is zero we need an alarm

The order system have a table with all received orders, but the requirement is that we need to monitor how many that has been received during the last 10 minutes. If this value is zero an alarm must be generated since that is an indication that the sales system is not generating orders. To achieve this monitoring we use the last and the previous sample of the total number of orders from example 1 and create the difference between the two to get a new virtual entity to measure, with a threshold level of 0.

Hopefully these examples have convinced you that with Bischeck you can do complex application monitoring by advanced threshold management and processes dependencies. And this can be done by just integrate it with your existing IT surveillance infrastructure.

But complex threshold management is not limited to the area of BAM. Even in the normal infrastructure and network monitoring the capabilities are important. A typical threshold for disk utilization is defined like 90% of total space. This is of course important but it could be equally important to understand the growth speed of the disk utilization, like if the size has grown with 5% from yesterday. This is simple in Bischeck without developing a specialized check command.

¹Nagios Core, Nagios XI, Icinga, OP5 Monitor, Shinken, Opsview, Centreon, NetEye, etc.

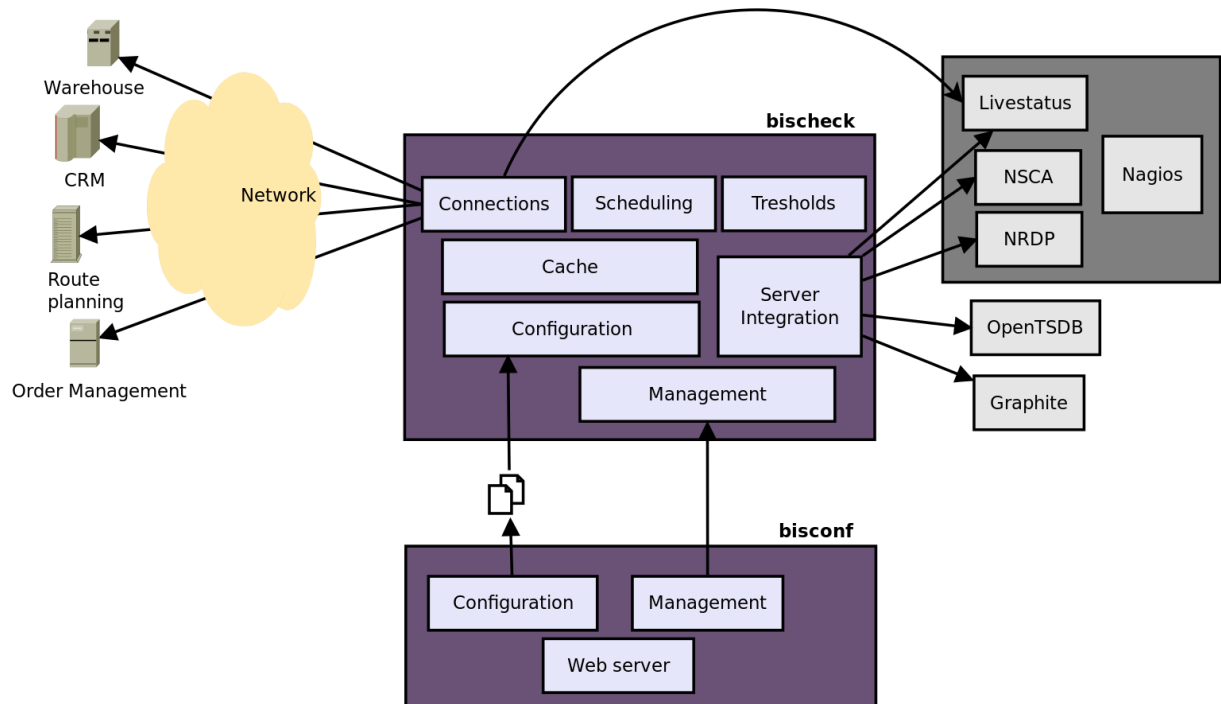


Figure 1.1: Architecture overview

The Bischeck architecture include the following main components:

- Configuration - Manage all configuration of Bischeck
- Scheduling - Controls the execution time of each service
- Connection - Manage connection management against the source systems to monitor
- Threshold - Manage the threshold configuration and processing
- Cache - Storage of all values that has been collected from the source systems
- Server integration - Manage integration and formatting to the different monitoring systems

Bisconf is an optional web tool that provide editing management of the Bischeck configuration files. In addition it also provide life cycle management to deploy new configuration and start and stop the Bischeck process.

Bischeck is open source and licensed under GPL version 2. Bischeck is written in Java and runs as a stand-alone daemon and integrates with different monitoring and surveillance tools over open protocols. If you have ideas to new features, find bugs, etc please visit <http://www.bischeck.org> where you can fill in bug reports and feature requests. You can also email anders.haal@ingby.com. We look forward to your feedback.

2 Service configuration

The design of Bischeck is based on the concept of configuring hosts, services and service items. As an example we have host *erpserver* and we define a service called *orders*. For the *order* service we can have multiple service items, defined like *ediOrders* and *mailOrders*. The service define a connection url to connect to the entity to measure and the service items define the statement to be executed to retrieve the entity value. The service url is used by Bischeck to instantiate the right service class to manage the the connection. For example if the url has a schema part that is jdbc, the JDBCService class is used. The mapping between the service url and the class to use is configurable. Each service class must implement the Service interface. The execution statement described for the service item must only return one single value. A service item class must implement the interface Serviceitem. The host can av multiple service, and each service can have multiple service items.

All host, service and service items are described in xml configuration files located in the *etc* directory of the installation, *\$BISHOME*. For more information about services and service items classes please section3.

3 Services and service items classes

The key configuration objects in Bischeck are the service and service item. The service key attributes are the service name, the schedules when to execute the service and the url that define the connection used by the service item.

The service item key attributes are the execute statement that define what should be executed to retrieve the measured entity and the threshold class used to process the measured entity to validate the state of the measured entity.

3.1 Service classes

The following Service classes is provided in the standard Bischeck distribution.

3.1.1 JDBC

JDBCService manage services connecting to databases over JDBC. JDBC driver jars should be put in the directory *customlib* in the Bischeck install directory to automatically be found at start-up.

A JDBC connection is described with a normal jdbc url and with an additional driver class in the driver tag of the Bischeck configuration.

3.1.2 Livestatus

LivestatusService enable connectivity to any Nagios server supporting MK Livestatus¹. With this service both the service state and performance data can be retrieved from a Nagios server. [New in 0.4.0]. A livestatus connection is specified with a url with the following format:

```
liveness: //hostname:port
```

The requirement for this type of connection is that livestatus is configured as a xinetd service on the server running livestatus and Nagios. To configure livestatus this way please see “Remote access to Livestatus via SSH or xinetd”.

3.1.3 Bischeck cache

LastCacheService enables connection to Bischeck’s internal cache where the measured data is stored. With the service any existing host->service->serviceitem can be retrieved from the cache. This can typical be used to create new virtual services based on

¹More information on MK Livestatus can be found on http://mathias-kettner.de/checkmk_livestatus.html

mathematical expression for combination of cache entities. The connection url for the Bischeck cache is:

```
bischeck://cache
```

3.1.4 Shell

[New in 0.4.1] *ShellService* enables execution of programs and scripts on the localhost. The connection url for the shell service is:

```
shell://localhost
```

3.2 Service item classes

The following *ServiceItem* classes is provided in the standard Bischeck distribution.

3.2.1 SQL

The *SQLServiceItem* class enables execution of a SQL statements. Its important that the SQL statement only return one single number value, like a "*select count(*)*" or "*select sum(orderValue) from*".

The *SQLServiceItem* class can be used when the service url is a *jdbc://...* url.

3.2.2 Livestatus

With the *LivestatusServiceItem* class both state and performance data can be queried for a Nagios service. The query is formatted as a json structure. To query a specific Nagios service state the following statement example can be used:

```
{"host ":" linux-server1 "," service ":" DNS "," query ":" state "}
```

The above will query the current Nagios state of service *DNS* for host *linux-server1*.

To query the the performance data for the Nagios service *HTTP* just change the query key value to *perfdata* and add the label key so the right label in a multi value performance data is retrieved:

```
{"host ":" linux-server1 "," service ":" HTTP "," query ":" perfdata "," label ":" size "}
```

The *LivestatusServiceItem* class can be used for service url *livestatus://hostname:port*.

3.2.3 Cache calculations

CalculateOnCache service item class enable statement execution using measured values residing in the Bischeck internal cache. Mathematical statements according to jep expressions capability is possible like:

```
if (( erpserver-orders-ediOrders[0] - erpserver-orders-ediOrders
    [1]) < 0,0, erpserver-orders-ediOrders[0] - erpserver-orders
    -ediOrders[1])
```

The above example check if the result of "erpserver-orders-ediOrders[0] - erpserver-orders-ediOrders[1]" is less then 0. If true, return 0 and if false return the result of "erpserver-orders-ediOrders[0] - erpserver-orders-ediOrders[1]". The value to use from the cache is defined by an index [X], where 0 is the latest value retrieved for the specific host-service-item.

The *CalculateOnCache* class can be used with service url *bischeck://cache*.

3.2.4 Check commands

[New in 0.4.1] With the *CheckCommandServiceItem* class performance data can be retrieved by executing a Nagios check command. The check command that is executed must output performance data according the Nagios specifications for the *CheckCommandServiceItem* to work. The execstatement query is formatted as a json structure. To execute a check command and retrieve a specific performance data the following statement example can be used:

```
{"check": "usr/lib/nagios/plugins/check_tcp -H localhost -p 22", "label": "time"}
```

The above will execute the check_tcp program and retrieve the performance data with label time.

The benefit of using a check command from Bischeck would be if there is a need to manage the performance data with dynamic thresholds that is not possible in Nagios Core.

Important is that there is no management of environment variables through Bischeck when executing a check command.

The *CheckCommandServiceItem* class can be used for service url *shell://localhost*.

4 Date formatting

Many service items will typically execute something depending on a date. For this reason Bischeck supports date macros in the execute statement for service items. For example, if this is a SQL select statement that using a date condition, the formatting could be done like this:

```
select count(orders) from order where  
fromdate='%%yyyy-MM-dd%%' and  
todate='%%yy.MM.dd%[D2]%%'
```

Bischeck will replace anything between %% and %% with the current date according to the format string. The formatting follows the structure of the Java SimpleDateFormat class. If the format string includes a %[macro, the current date will be calculated, where Y means year, M means month and D means day. So %[D-1] means subtract one day from the current date and %[M2] means add two months to the current date. A construction like D-1Y-1 is not supported.

5 Service item cache - LastStatusCache

Bischeck keeps a cache of the measured data points for every service items. The cache content can be used in the specification of execute statements of serviceitems and in threshold classes to create dynamic and adaptive thresholds, see section 6.1 on page 18.

The cached data is saved to disk every time Bischeck is stopped and reloaded every time Bischeck is started. The stored data is written to a file called *lastStatusCacheDump* and located in the directory specified by the property *lastStatusCacheDumpDir*, see 8.3. Before a save is made a backup file is created in the same directory called *lastStatusCacheDump.bak*.

5.1 Methods to retrieve cache data

There is a number of ways to retrieve data from the cache. The cache support both retrieval by index or by time, and the retrieval can be done for a single element or to a list of elements.

5.1.1 Retrieve single cache data element by index

For each service item the last 500 measured values are kept in the cache in a LRU list. To retrieve data from the cache the following format is used, *host-service-item[index]* where index is the lru entry, with 0 specifying that last retrieved measured value , e.g. *erpserver-orders-ediOrders[0]*.

5.1.2 Retrieve single cache data element by time

Data can also be retrieved from the cache by using a time offset, like the cache data of *erpserver-orders-ediOrders* 30 minutes ago. This is specified by setting the index part to *-30M*, *erpserver-orders-ediOrders[-30M]*. This is useful since the relation between static index and service schedule do not need to be known. The time operator supported is S for seconds, M for minutes and H for hours.

5.1.3 Retrieve a list of cache data elements by index

For mathematical functions that operates of a list of elements, like a sum, avg, max, Bischeck support two methods to retrieve the range by index. The first method is to define a from “to index”, like *erpserver-orders-ediOrders[0:9]* that will retrieve all cache elements from index 0 to 9 as a list separated by “,”¹.

¹The choice of , as separator character is due to this is the separator character used by JEP functions that take a list of values.

```
sum(erphost-orders-ediorders[0:9])
```

This will calculate the sum of the last 10 cached elements for *erpserver-orders-ediOrders*.

The second methods is to define a list of index that do not need to be in sequence, like *erpserver-orders-ediOrders[1,3,5]* that will retrieve the elements for index 1, 3 and 5 as a list separated by “,”.

```
sum(erphost-orders-ediorders[1,3,5])
```

5.1.4 Retrieve a list of cache data elements by time range

To retrieve list of the data elements based on a time range the expression has the format *erpserver-orders-ediOrders[-30M:-120M]*. This will retrieve a list of all elements that exist in the cache between 30 minutes to 120 minutes ago. The time operator supported is S for seconds, M for minutes and H for hours.

5.2 Mathematical expressions on cache data

In places where cache data can be used, like thresholds and serviceitems, it's possible to use mathematical expressions on the data. In Bischek we use the mathematical package JEP to evaluate mathematical expressions and formulas. Lets look at some examples.

```
erphost-orders-ediorders[0] / erphost-orders-ediorders[-30M]
```

This example will divide the last cached data, index 0, for *erphost-orders-ediorders* with the cached data of *erphost-orders-ediorders* 30 minutes ago.

```
max(erphost-orders-ediorders[-5M], erphost-orders-ediorders[-10M],  
    erphost-orders-ediorders[-15M])
```

In the above the highest value will be returned of the cached data for *erphost-orders-ediorders* cached 5, 10 and 15 minutes ago.

For more information of all functions that can be used please visit www.sourceforge.net and download the JEP distribution to get access to the JEP version 2 documentation.

5.2.1 JEP standard functions

The standard JEP functions supported are:

- round(x,y) - rounding the value of x to y number of decimals.
- ceil(x) - calculate the ceiling value of x.
- floor(x) - calculate the floor value of x.
- abs(x) - calculate the absolute value of x.
- mod(x,y) - calculate the modulus of x % y.

- `sqrt(x)` - calculate the square root of x.
- `sum(x1,x2,x3,...)` - calculate the sum of the list of values.
- `ln(x)` - calculate the natural logarithm of x.
- `log(x)` - calculate the logarithm base 10 of x.
- `exp(x)` - calculate the exponential value of x.
- `pow(x,y)` - calculate the x as the power of y.
- `rand()` - random value between 0 and 1
- `if(cond, trueval, falseval)` - if the cond expression is true trueval is returned, if false falseval is returned.

5.2.2 JEP extensions

With JEP it's easy to create additional functions. With Bischek the following functions has been added as part of the distribution.

- `avg(x1,x2,x3,...)` - Calculate the average of the list of numbers.
- `max(x1,x2,x3,...)` - Calculate the max value of the list of numbers.
- `min(x1,x2,x3,...)` - Calculate the min value of the list of numbers.
- `multNull((x1,x2,x3,...)` - Multiply the numbers in the list, but if any of the numbers are null the function will return null.
- `divNull(x1,x2)` - Divide the numbers, but if any of the numbers are null the function will return null.

5.3 Null value in the cache

Null values represent the existent of no data. A null value in the cache can exists due to different reasons.

1. The connection specified in the service fails. By default a null value will not be stored when this happens but if the property *saveNullOnConnectionError* is set to true a null value will be inserted in the cache.
2. If the execstatement specified for the serviceitem return null, a null value is inserted in the cache.

When accessing the cache by index or by time like, `erphost-orders-ediorders[10]` or `erphost-orders-ediorders[-5M]` , it will return null:

- If null is stored in the cache at the index.

- If null is stored in the cache at the index that the time specification resolves to.
- If the index is not yet been populated with data.
- If the index is out of bound of the specified cache size.

When using cache data with mathematical function that takes a list of arguments, the null value can be handled in different ways depending how the cache data is retrieved and how the property *notFullListParse* is set. Lets show an example to explain the different scenarios.

```
sum(erphost-orders-ediorders[0] , erphost-orders-phoneorders[0])
  * erphost-orders-state[0]
```

In the above example the *sum* function will evaluate to null if *erphost-orders-ediorders[0]* and/or *erphost-orders-phoneorders[0]* return null. If the the property *notFullListParse* is set the *sum* function will only return null if both *erphost-orders-ediorders[0]* **and** *erphost-orders-phoneorders[0]* return null. The whole statement will return null independent of the *notFullListParse* property if *erphost-orders-state[0]* returns null.

By the example we can see that the property *notFullListParse* has effect on how null is managed for functions that manage a list of data like *sum*, *avg*, *max*, *min*, etc. What the *notFullListParse* property do is that it will filter out all null values when they occur in a function that take a list of data. This works in the same way if ranges are used as the example below.

```
sum(erphost-orders-ediorders[0:100] , erphost-orders-phoneorders
    [-5M:-15M]) * erphost-orders-state[0]
```

If the *notFullListParse* is set to true all the cache data in *erphost-orders-ediorders[0:100]* and *erphost-orders-phoneorders[-5M:-15M]* will be filtered out. For functions like *avg* the average will only be based on the none null values.

The only exception to this rule is when using time range like *erphost-orders-phoneorders[-5M:-15M]*. A time range will resolve to null if the lower or higher boundaries are not in the range of the data that exists in the cache.

The *notFullListParse* property also have another implication and that if using mathematical expressions inside a function that take a list of values.

```
avg(erphost-orders-ediorders[0] * 2, erphost-orders-phoneorders
    [0] * 3)
```

The above example will result in null for the whole *avg* function if any of *erphost-orders-ediorders[0]* or *erphost-orders-phoneorders[0]* is null since the multiply with a constant will fail if any of the cache data are null. To manage this the *multNull*² function must be used instead of the *** operator.

```
avg(multNull(erphost-orders-ediorders[0],2) , multNull(erphost-
    orders-phoneorders[0],3))
```

²See 5.2.2 for more information about functions that manage null.

Using multNull will result in that the *avg* function will return null only if both erphost-orders-ediorders[0] and erphost-orders-phoneorders[0] are null.

6 Threshold configuration

To define the state of the retrieved service items value a threshold class is specified for each service item. The threshold class evaluate the measured value with some logic and return the state as okay, warning or critical level. The threshold instance for the specific service item is stored in a threshold cache and valid for a period of a day. Every new day the threshold cache is invalidated and Bischeck looks for new valid threshold classes to instantiate for the service item to be used for the period of a day. The reason that the period is set to a day is that there may be a need to configure the threshold instance differently depending on the day of the week or month.

A valid threshold class must implement the interface `Threshold`. This structure enables a flexible implementation of very different ways to calculate or specify a threshold. In the simplest form a threshold class could just return a okay independent of measured value and in the more complex solution be based on algorithms, database content, measured values from other service items, day of month or some complex combination.

Warning and critical level specification are also part of the threshold class. The threshold class are also responsible to define what operations are support for the measured value, like measured value should be higher, lower or in an interval of the threshold. Warning and critical level are always defined as the percentage of the threshold.

6.1 Twenty4HourThreshold class

The `Twenty4Hour` threshold class divides the day into 24 hours. For each hour of the day a threshold is defined. The two threshold values that are next to each other used to calculate a slope of a linear equation between the two closest hours. For example, if the threshold value is 1000 at 14:00 and 1600 at 15:00, the calculation for a threshold value between 14:00 and 15:00 is $y=x*(1600-1000)/60 + 1000$. At 14:20 the threshold is $20*(1600-1000)/60+1000 = 1200$.

The threshold model gives a linear equation with one hour granularity, but over 24 hours it can resemble a curve. This behavior is typical in business systems where the key business values are distributed in a none linear and none constant way over the period of a day, e.g. incoming orders.

6.1.1 Period definition

Since the threshold for a service item can be different depending on the month, day of month, week, day of week the configuration supports thresholds to be described on a granularity called period. A period include multiple months and weeks definitions as long as they share the same threshold definition. For a month it is possible to specify

a specific month and/or a day of a month and for a week a specific week and/or a day of a week. To find the right threshold period the the class look for threshold period specification in the following order:

1. Month and day of month
2. Week and day of week
3. Day in month
4. Day in the week
5. Month
6. Week
7. Default

Month is specified between 1-12, week 1-53, day of month 1-31 and day of week 1-7 (1=Sunday). The default threshold period is used if no other matching occurs.

Since holidays are often days where the business are not operational there is a way to describe days that should no have any threshold checks. These exclude days are checked before any other rule above are evaluated.

6.1.2 Calculation definition

The class support tree ways how the threshold is compered to the measured value:

">" Measured value should be higher then threshold. If the measured value is lower then $\text{threshold} * \text{warning}(\%)$ warning state is set and if measured value is lower then $\text{threshold} * \text{critical}(\%)$ critical state is set.

"<" Measured value should be lower then threshold. If the measured value is higher then $\text{threshold} * (1 - \text{warning}(\%))$ warning state is set and if if measured value is higher then $\text{threshold} * (1 - \text{critical}(\%))$ critical state is set.

"=" Measured should be in the interval of the threshold. If the measured value is lower then $\text{threshold} * \text{warning}(\%)$ OR higher then $\text{threshold} * (1 - \text{warning}(\%))$ a warning state is set and if measured value is lower then $\text{threshold} * \text{critical}(\%)$ OR higher then $\text{threshold} * (1 - \text{critical}(\%))$ a critical state is set.

For complete configuration description see section 8.6 on page 31.

6.1.3 Hours and threshold definition

The measured value is compered against the threshold value that is calculated from linear equation of the two closest threshold values. As described in the introduction of this chapter we can set the threshold values to fixed number, but in a business system this is not enough. Let take an example. The number of orders that can be invoiced

during a day is probably depending on the number of orders received. So instead of setting the threshold to a fixed number we can use an expression based threshold like “80 % of received orders”. The syntax of expression based thresholds is simple and powerful. Expressions are based on the JEP package, see chapter 16, where the parameters are any measured values that exist in the service item cache, see 5. For example the expression *"erpserver-orders-ediorders[0]*0.8"* is the threshold is set to 80% of the last measured value of the service item *ediorders* for the service *order* and host *erpserver*. We could also combine multiple cached values from different sources in the same expression like, *"erpserver-orders-ediorders[0] / geoserver-route-finalroute[0]"* to get some sort of ratio threshold.

If we just need to check parts of a day for thresholds, just set the hour to NULL and no calculation will be done for that time interval. For example if no threshold checks should be done on weekends just define period rule number 4 for day 1 and 6 with all hour thresholds set to NULL.

7 Server integration

Since Bisccheck support integration with multiple monitoring infrastructures its not possible to define a single message format or protocol. For that reason each Server class implementation is responsible for message formatting and communication with the integrated server. If the server system is not responding the message will be dropped on the Bisccheck side.

7.1 Nagios integration with NSCA, NRDP and Livestatus

The data format between Bisccheck and Nagios follow the standard Nagios format. This means that the Nagios host and service names must be the same as you configure in Bisccheck. Since NSCA and NRDP [\[New in 0.4.1\]](#) works with passive checks the *check_dummy* or equivalent check command should be used on the Nagios side in the service configuration. For a service in Bisccheck the data is split in the status and the performance output. The Bisccheck status output has the following format:

```
<level> <service item name> = <measured value> (<threshold> <warning value>
<warning_calc method> <critical value>) <critical calc method> , <service item name>
= ...
```

Since a service can have multiple service items, the output is presented as a concatenated string of the service items. The *<measured value>* is the value that was retrieved from the execution of the specific service item execute statement. The *<threshold>* is the current threshold value that the measured value has been compared against. The calculated warning and critical levels are also calculated based on the percentage value of the threshold.

Depending of the calculation method the string representation will differ. The following methods are support:

- Measure value must be higher the threshold: *<warning level> = > W > <critical level> = > C >*
- Measure value must be lower the threshold: *<warning level> = < W < <critical level> = < C <*
- Measure value must be in the interval to the threshold: *<warning level> = = +-W = <critical level> = = +-C =*

Example of the output:

```
OK ediOrders = 12000 (11000 > W > 9900 > C > 7700)
```

In this case the threshold value is 11000 and warning and critical levels has been set to 10% and 30 % of the current threshold. Since the warning and critical level are as a fixed percentage the calculated level will changed with the current threshold.

If there are no threshold values defined for the current period in which the measured value is done, the threshold is reported as null and no calculation is done that can be used for notification. Null can also be reported as the measured value. This can occur if a none or a null value is retrieved for the service item, for example from a faulty SQL statement.

Service connection problems will be reported as critical. Warning and critical notification will be based on the service item that has the highest level of severity if there are more then one service item defined for a service. If one service item reports critical and the other reports OK the service will report critical.

For the performance data, which Nagios graphs, like pnp4nagios, are based upon, each service item is included with the addition of the current threshold.

```
ediOrders=12000;9900;7000;0; threshold=11000;0;0;0; avg-exec-time=223ms
```

The execution time of the service execute statement is always part of the performance data and reported in milliseconds. With the pnp4nagios template that is provided by the Bischeck installation the average execution time is not graphed.

7.1.1 NSCA configuration

The configuration parameters that need to be define in the server.xml for NSCA are:

- class - NSCAServer
- hostAddress - the IP or name for the NSCA server, default is localhost.
- port - the socket port number on the NSCA server where NSCA is listening, default is 5667.
- connectionTimeout - the timeout time in milliseconds before Bischeck give up on the the connection, default is 5000.
- encryptionMode - the encryption mode used in the communication with the NSCA server. Supported encryption modes are NONE, XOR, and TRIPLE_DES, default is XOR.
- password - the NSCA server password, no by default.

7.1.2 NRDP configuration

The configuration parameters that need to be define in the server.xml for NRDP are:

- class - NRDPSTServer
- hostAddress - the IP or name for the NRDP server, default is localhost.

- port - the socket port number on the NRDP server where NRDP is listening, default is 80.
- connectionTimeout - the timeout time in milliseconds before Bischeck give up on the the connection, default is 5000.
- path - the URL path for NRDP, default is nrdp.
- password - the NRDP server token, no by default.

7.1.3 Livestatus configuration

The configuration parameters that need to be define in the server.xml for Livestatus are:

- class - LiveStatusServer
- hostAddress - the IP or name for the Livestatus server, default is localhost.
- port - the socket port number on the server where Livestatus is listening, default is 6557.
- connectionTimeout - the timeout time in milliseconds before Bischeck give up on the the connection, default is 5000.

7.1.4 Nagios state on null

If the service connection fail or the service item execstatement returns a null value its not possible to calculate thresholds or to determine the state that should be propagated to Nagios. The state that should be communicate to Nagios in this type of situation is defined by property *stateOnNull*, that can be set to CRITICAL (2), WARNING (1), OK (0) or UNKNOWN (3). Default is UNKNOWN.

7.2 OpenTSDB

OpenTSDB is an monitoring system that provide storage and indexing of high volume of time series data. Access to the data in through a graphical web console provide easy way to view many different data sources for advanced analysis. The OpenTSDB integration is provided over OpenTSDB text based protocol. From Bischeck the format will be the following for the host erpserver, service orders and service item ediOrders:

```
put bischeck.measured 1288946927 12000 host=erpserver service=
  orders serviceitem=ediOrders
put bischeck.threshold 1288946927 11000 host=erpserver service=
  orders serviceitem=ediOrders
put bischeck.warning 1288946927 9900 host=erpserver service=
  orders serviceitem=ediOrders
put bischeck.critical 1288946927 7700 host=erpserver service=
  orders serviceitem=ediOrders
```

So for each host, service and service item configuration 4 lines of data will be sent to the OpenTSDB server for the measured, threshold, warning and critical value.

The configuration parameters that need to be define in the server.xml for OpenTSDB are:

- class - OpenTSDBServer
- hostAddress - the IP or name for the OpenTSDB server, default is localhost.
- port - the socket port number on the OpenTSDB server where OpenTSDB is listening, default is 4242.
- connectionTimeout - the timeout time in milliseconds before Bischek give up on the the connection, default is 5000.

7.3 Graphite

Graphite is an monitoring system that provide excellent graphing. Access to the data in through a graphical web console provide easy way to view many different data sources for advanced analysis. The Graphite integration is provided over Graphites text based protocol. The format has the following simple format:

```
metric_path value timestamp\n
```

The metric_path is the dot separated string like foo.bar.baz. Value is a number and timestamp is the UNIX epoch time in milliseconds.

From Bischek the format will be the following for the host erpservers, service orders and service item ediOrders:

```
erpservers.orders.ediOrders.measured 12000 1288946927\n
erpservers.orders.ediOrders.threshold 11000 1288946927\n
erpservers.orders.ediOrders.warning 9900 1288946927\n
erpservers.orders.ediOrders.critical 7700 1288946927\n
```

So for each host, service and service item configuration 4 lines of data will be sent to the Graphite server for the measured, threshold, warning and critical value.

The Graphite server integration support a method to filter what data that will be sent to Graphite. The filter is defined by regular expression with the property *doNotSendRegex*. If a host-service-serviceitem in anyway match the regular expression it will not be sent to Graphite. Multiple regular expressions can be defined by a list of regular expressions in the property *doNotSendRegex* by setting a delimiter between the regular expressions defined by the property *doNotSendRegexDelim*. For example the regular expression *^erpservers* will not send any host-service-serviceitem that begins with the string erpservers.

The configuration parameters that need to be define in the server.xml for OpenTSDB are:

- class - GraphiteServer
- hostAddress - the IP or name for the Graphite server, default is localhost.
- port - the socket port number on the Graphite server where Graphite is listening, default is 2003.
- connectionTimeout - the timeout time in milliseconds before Bischeck give up on the the connection, default is 5000.
- doNotSendRegex - A list of regular expressions. Default is no regular expressions.
- doNotSendRegexDelim - The delimiter to separate regular expressions in the property doNotSendRegex, default is "%".

8 Bischack configuration files

The basic configuration file are xml based and located in the *\$BISHOME/etc* directory. The distribution package also include all xml schema, xsd, files used for validation. For detail information please review the xsd files located in the directory *\$BISHOME/resources*. Remember that all xml configuration files should use HTML encoded characters.

8.1 Naming standard

The rules for naming hosts, services and serviceitems must adhere to the following rules:

- hostname - can include any character from a-z, A-Z and 0-9. The hostname can also include the following special character, dash (-), dot (.) and underscore (_). Its not allowed to start the name with any of the special characters. Examples of valid names are erphost.xyz.com and erp1_host.my-domain.com.
- service and serviceitem - can include any character from a-z, A-Z and 0-9. The name can also include the following special characters, dash (-), dot (.), underscore (_), at sign (@) and space (). Its not allowed to start or end the name with dash (-) and space ().

When using a host-servie-serviceitem[X] expression in the execstatement tag of a serviceitem in the bischeck.xml configuration file or in the hour tag of the 24threshold.xml configuration file, a dash (-) in any of the names must be quoted with a backslash (\), like:

```
erp1_host.my\-domain.com-@service1-item.1.1[1]
```

8.2 bischeck.xml

The Bischack configuration can be described as a hierarchy of hosts, services and service items to monitor. Each host can have one or more services, and for each service one or more service items can be configured.

```
<bischeck> 1
  <host> 2
    <name>erpserver</name> 3
    <desc>ERP server</desc> 4
    <service> 5
      <name>orders</name> 6
```

```

<desc>Order management</desc> 7
<schedule>0 0/5 * * * ?</schedule> 8
<url> 9
    jdbc:mysql://erphost/erpdb?user=bischeck&password 10
        =bischeck
</url> 11
<driver> 12
    com.mysql.jdbc.Driver 13
</driver> 14
<serviceitem> 15
    <name>ediorders</name> 16
    <desc>Inbound edi orders</desc> 17
    <execstatement> 18
        select count(*) from orders where createdate=&apos; 19
            ;%%yyyy-MM-dd%%&apos;;
    </execstatement> 20
    <thresholdclass> 21
        Twenty4HourThreshold 22
    </thresholdclass> 23
    <serviceitemclass> 24
        SQLServiceItem 25
    </serviceitemclass> 26
</serviceitem> 27
</service> 28
</host> 29
</bischeck> 30

```

In the host section the following elements are defined:

- name - the name of the host. ¹
- desc - a optional description field

In the service section the following elements are defined:

- name - name of the service. ²
- desc - a optional description field
- schedule - one to many service execution schedules can be defined. See the section 9 for more information about configuration options.
- sendserver - true or false to define if the service data should be sent to the configured monitoring servers, see 8.5. [New in 0.4.0]. If set to false the measured data will only be placed in the Bischeck cache and be used to create virtual services and threshold rules.

¹Must be the same name as the host name is configured with on the Nagios server.

²Must be the same name as the host name is configured with on the Nagios server.

- url - the connection specification for the server/service to monitor in a url format
- driver - a class that is needed by the url specification. The driver class should be placed in the *customlib* directory.

In the service item section the following elements are defined:

- name - name of the service item
- desc - a optional description field
- execstatement - the specification of what to be executed by the service item. The statement will be dependent on the type of service item class used.
- serviceitemclass - the ServiceItem class to use for the service item. For ServiceItem classes that are part of the distribution the class name is enough. For custom developed ServiceItem classes the full class name should be specified and the classes or jar files should be located in the *customlib* directory.
- thresholdclass - the Threshold class to be used by the service item configuration. For Threshold classes part of the distribution the class name is enough. For custom developed Threshold classes the full class name should be specified and the classes or jar files should be located in the *customlib* directory.

8.3 properties.xml

The properties.xml include properties used by the core of Bischeck, but can also be used by custom developed services and threshold classes. The properties xml has a simple structure of key/value pair:

```
<properties> 1
  <property> 2
    <key>akey</key> 3
    <value>avalue</value> 4
  </property> 5
</properties> 6
```

Any properties should have a default value that is set in code. The following properties are currently used by core Bischeck:

- thresholdCacheClear- the time when the threshold cache should be evicted each day, default is "10 0 00 * * ? *".
- pidfile - the pid file for Bischeck, default is *"`/var/tmp/bischeck.pid`"*.
- mindaysinfirstweek - define the week numbering according to ISO 8601, default is 4.
- firstdayofweek - define the week numbering according to ISO 8601, default is 2.

- `lastStatusCacheDumpDir` - the directory where cache is stored, default is `/var/tmp`.
- `lastStatusCacheSize` - the maximum size of cached data per host, service and serviceitem, default is 500.
- `runAfterDelay` - the delay time in seconds used when a service is scheduled to run with “Service relation scheduling”, see 9. Default is 10 seconds.
- `notFullListParse` - if set to true a reference to a cache item where a list is used, like `erpserver-orders-ediorders[1:10]`, will return data even if not all index in the list has data. Read more about how this property affect the retrieval of cache data in section 5.3.
- `stateOnNull` - the Nagios state to be used if a execstatment in a service item return null. The property can be set to an integer 0,1,2 or 3 or to a string OK, WARNING, CRITICAL or UNKNOWN. The default is UNKNOWN.
- `saveNullOnConnectionError` - this property define if a null value should be stored in the cache if you have your service class has a connection exception of any sort. The default is false.

The `properties.xml` file can also include class specific properties typical used for custom developed service, service item and threshold classes. A class specific property should have a key formatted in the following way - `classname.propertyname`

Any class specific properties must have a default value implemented by the class itself. E.g. for `JDBCService` there is a property called `querytimeout` that sets the max time in seconds before aborting the query. This property has the name `JDBCService.querytimeout` and have a default value of 10 seconds.

8.4 urlservice.xml

The configuration define of the mapping between service url schema and Service class. The `urlservice.xml` has a structure of:

```

<urlservices>
  <urlproperty>
    <key>jdbc</key>
    <value>JDBCService</value>
  </urlproperty>
  <urlproperty>
    <key>bischeck</key>
    <value>LastCacheService</value>
  </urlproperty>
</urlservices>

```

- `key` – the schema part of the service url, i.e. `jdbc`.

- value – the Service class name. For Service classes part of the distribution the class name is enough to specify. For custom developed Service classes the full class name should be specified and the classes or jar files should be located in the *customlib* directory.

8.5 server.xml

The configuration file define the mapping between service url schema and Server class. The urlservice.xml has a structure of:

```

<servers>
  <server name="Nagios1">
    <class>NSCASServer</class>
    <property>
      <key>hostAddress</key>
      <value>localhost</value>
    </property>
    <property>
      <key>encryptionMode</key>
      <value>XOR</value>
    </property>
    <property>
      <key>password</key>
      <value>change this</value>
    </property>
    <property>
      <key>port</key>
      <value>5667</value>
    </property>
    <property>
      <key>connectionTimeout</key>
      <value>5000</value>
    </property>
  </server>

  <server name="Nagios2">
    .....
  </server>
</servers>

```

Each server that is integrated with Bischeck must be defined with the server tag and with a unique attribute name. This solution enable Bischeck to send data to multiple servers of the same type, like multiple Nagios servers. The class tag define what Server class implementation to use for the integration. Each Server implementation can have

different numbers of properties that is specific for the Server class. Custom developed server classes should be placed in the *customlib* directory

8.6 24thresholds.xml

The threshold class Twenty4HourThreshold described in section 6.1 is configured through the 24thresholds.xml file.

```

<twenty4threshold>
  <servicedef>
    <hostname>erpserver</hostname>
    <servicename>shipments</servicename>
    <serviceitemname>outboundshipment</serviceitemname>

    <period>
      <!-- valid for any 21th day in the month -->
      <months>
        <dayofmonth>21</dayofmonth>
      </months>
      <!-- valid for week 12 (middle of March)
           and if its a Thursday -->
      <weeks>
        <week>12</week>
        <dayofweek>5</dayofweek>
      </weeks>

      <calcmethod>&gt;</calcmethod>
      <warning>10</warning>
      <critical>30</critical>
      <hoursIDREF>1</hoursIDREF>
    </period>

    <period>
      <!-- valid if its a Friday -->
      <weeks>
        <dayofweek>6</dayofweek>
      </weeks>
      <calcmethod>&gt;</calcmethod>
      <warning>10</warning>
      <critical>30</critical>
      <hoursIDREF>2</hoursIDREF>
    </period>

    <period>

```

```

        <!-- This will be used if no                                37
            other rule if applicable -->                             38
        <calcmethod>&gt;</calcmethod>                                   39
        <warning>10</warning>                                       40
        <critical>30</critical>                                     41
        <hoursIDREF>31</hoursIDREF>                                 42
    </period>                                                       43
                                                                    44
</servicedef>                                                       45
                                                                    46
<hours hoursID="1">                                                 47
    <!-- 00:00 -->                                                  48
    <hour>500</hour>                                                49
    <!-- 01:00 -->                                                  50
    <hour>1500</hour>                                              51
    <!-- 02:00 -->                                                  52
    <hour>4000</hour>                                              53
    .....                                                         54
    <!-- 21:00 -->                                                  55
    <hour>9000</hour>                                              56
    <!-- 22:00 -->                                                  57
    <hour>10000</hour>                                           58
    <!-- 23:00 -->                                                  59
    <hour>11000</hour>                                           60
</hours>                                                            61
                                                                    62
<hours hoursID="2">                                                 63
    <!-- 00:00 -->                                                  64
    <hour>1500</hour>                                              65
    <!-- 01:00 -->                                                  66
    <hour>2500</hour>                                              67
    <!-- 02:00 -->                                                  68
    <hour>5000</hour>                                              69
    .....                                                         70
    <!-- 21:00 -->                                                  71
    <hour>10000</hour>                                           72
    <!-- 22:00 -->                                                  73
    <hour>12000</hour>                                           74
    <!-- 23:00 -->                                                  75
    <hour>14000</hour>                                           76
</hours>                                                            77
                                                                    78
<hours hoursID="2">                                                 79
                                                                    80

```


<hourinterval>	81
<from>09:00</from>	82
<to>12:00</to>	83
<threshold>erpserver-orders-ediorders[0]*0.8</hour>	84
</hourinterval>	85
	86
<hourinterval>	87
<from>12:00</from>	88
<to>15:00</to>	89
<threshold>erpserver-orders-ediorders[0]*0.4</hour>	90
</hourinterval>	91
	92
<hourinterval>	93
<from>16:00</from>	94
<to>17:00</to>	95
<threshold>erpserver-orders-ediorders[0]*0.2</hour>	96
</hourinterval>	97
	98
<hourinterval>	99
<from>20:00</from>	100
<to>22:00</to>	101
<threshold>10000</hour>	102
</hourinterval>	103
	104
	105
</hours>	106
	107
<!-- Holidays -->	108
<holiday year="2011">	109
<dayofyear>0101</dayofyear>	110
.....	111
<dayofyear>1224</dayofyear>	112
<dayofyear>1225</dayofyear>	113
</holiday>	114
	115
</twenty4threshold>	116

The configuration is based on two main parts, a service definition called *servicedef* tag and an *hours* tag. For each combination of host, service and serviceitem a *servicedef* tag is specified. Each *servicedef* can have one to many period specifications, see 6.1.1. The period specify what hoursID that should be used for a specific combination of months and weeks definitions, see 6.1.3for more information. The period also define calculation method, warning and critical level, see 6.1.2.

The *hours* tag support two different formats, the 24 hour listing or a from-to listing.

If the *hours* tag use the 24 hour listing there must be 24 *hour* tags, each represent on hour of the day. An hour tag can have a null value, meaning no threshold, a fixed value or an mathematical expression according to JEP, but with variables from any measured value existing in the last value cache, see 5. For more information about threshold specifications please see 6.

If the from-to format is used there can be one to many *hourintervals* that define what threshold that should be used between the from time to the to time. Only full hours are currently supported. The following rules are valid for this format:

- If there is overlapping hours exist the last *hourinterval* tag in the list have precedence. In the above listing that means that the threshold for 12:00 are “erpserver-orders-ediorders[0]*0.4”
- Any hour that are not defined are treated as to have no threshold. In the above example there are no thresholds for 00:00-09:00, 17:00-20:00 and 22:00-00:00.

The final section, holidays, describe days of the year where no threshold will be tested. For that reason the service will always return an OK state for a Bischeck services using the threshold class `Twenty4HourThreshold`.

9 Service scheduling

The service scheduling define when a service is triggered to be executed. For each service there can be multiple scheduling configurations, but at least one is mandatory. The scheduling can be described in three ways.

9.1 Interval scheduling

The simple format describe a interval execution that are repeated forever. The format is just a number and a indicator defining the granularity in seconds (S), minutes (M) or hours (H). 10M specify that the service should be executed every ten minutes.

```
<schedule>10M</schedule>
```

1

The initial start time for an interval based schedule is calculated as a time randomly in the specified interval. For 10M this means that the service first schedule is between 0 to 10 minutes from the start time of Bischek or at a reload.

9.2 Cron based scheduling

The second format is more advanced and follow the cron specification of Quartz, see <http://www.quartz-scheduler.org>. With this format its possible to define scheduling expression like “0 15 10 ? * MON-FRI” which would schedule the service at 10:15am every Monday, Tuesday, Wednesday, Thursday and Friday. For more cron examples please visit <http://www.quartz-scheduler.org/documentation/quartz-2.1.x/tutorials/crontrigger>

```
<schedule>0 15 10 ? * MON-FRI</schedule>
```

1

9.3 Service relation scheduling

[New in 0.4.0]The service relation scheduling trigger a service to execute after another service has been scheduled. This is useful when a service is depending on data for another service for its thresholds or execution statement. The format is “*host-service*” specifying the host and service that will trigger the execution of the service.

```
<schedule>erpserver - orders</schedule>
```

1

10 Configuration tools

Bischeck provides by default no additional tools for configuration other than the normal editor to manage the xml configuration files. The default configuration files in *etc* directory are valid for the test example configuration, see 13.2.

A web based configuration tool, bisconf, is under development. Bisconf is hosted on the Bischeck project site, <http://www.bischeck.org>.

11 Bischack internal surveillance

Bischack use the Java JMX standard for internal monitoring. Please read the javadoc for the classes `ExecuteMBean` and `LastStatusCacheMBean` to understand the methods available or just start *jconsole* or equivalent tool.

To allow remote JMX connections the following JMX settings are used by the *bischack* script located in the *\$BISHOME/bin* directory.

```
jmxport=-Dcom.sun.management.jmxremote.port=<set by the
  installer script>
jmxrmiserver=-Djava.rmi.server.hostname=<set by the installer
  script>
jmxauth=-Dcom.sun.management.jmxremote.authenticate=<set by
  the installer script>
jmxssl=-Dcom.sun.management.jmxremote.ssl=false
jmxpasswd=-Dcom.sun.management.jmxremote.password.file=
  $bishome/etc/jmxremote.password
jmxaccess=-Dcom.sun.management.jmxremote.access.file=$bishome
  /etc/jmxremote.access
```

If authentication should be used please change the configuration of the roles and passwords in the files `jmxremote.access` and `jmxremote.password`. For more information how to configure JMX please visit

<http://www.oracle.com/technetwork/java/javase/tech/javamanagement-140525.html>.

JMX is only enabled when the *bischack* script is called with the argument "*Execute -d*", which is the way the *bischackd* init script call the script *bischack* to start Bischack in daemon mode.

11.1 Bischack MBeans

The Bischack JMX MBeans provide instrumentation of the running Bischack daemon. From an operations perspective Bischack can be shutdown and reloaded through JMX. A reload operations is used to force Bischack to re-read it's configuration files but without restarting the process.

Different execution times of the Bischack daemon are also exposed through JMX. All timers are based on the <http://metrics.codahale.com/> and shows the following parameter for each defined timer:

- FifteenMinuteRate
- FiveMinuteRate

- OneMinuteRate
- MeanRate
- Count
- Max
- Min
- Mean
- Standard deviation (StdDev)
- 50th Percentile
- 75th Percentile
- 95th Percentile
- 98th Percentile
- 99th Percentile
- 999th Percentile

The timers used in Bischeck are:

- *com.ingby.socbox.bischeck.service:name=execute,type=ServiceJob*
This is the most important metric since it shows the total execution time to process a definition of a host, service and serviceitem(s) every time it's scheduled. The execution time included connecting and retrieving data, execution over serviceitem logic, threshold calculation and the time to send data to the defined servers.
- *com.ingby.socbox.bischeck.servers:name=execute,type=ServerExecutor*
This is the total time to send the data to all the servers defined. This time in is included in the *ServiceJob* time. Each server also have it's own timer to track the time down on individual servers.
- *com.ingby.socbox.bischeck.jepext:name=calulate,type=ExecuteJEP*
The execution time to calculate any mathematical expressions like sum, average, etc used in serviceitem and threshold logic.
- *com.ingby.socbox.bischeck.cache.provider:name=parse,type=LastStatusCache*
The execution time to parse and retrieve data from the cache used in serviceitem and threshold logic. This include all different way to retrieved data like based on index, index range, time and time range.

12 Building Bischeck

To build Bischeck from source is simple. Check out the Bischeck trunk from gforge.ingby.com:

```
$ svn checkout --username anonymous http://gforge.ingby.com/  
    svn/bischeck/trunk bischeck
```

To build a Bischeck distribution run from the directory where you checked out the Bischeck code:

```
$ ant dist
```

This will create a compressed tar file in the *target* directory, named `bischeck-x.y.z.tgz` where `x.y.z` is the version number. Different versions of Bischeck can be checked out from the tags directory located in <http://gforge.ingby.com/svn/bischeck/tags>

12.1 Jar customization

To support custom jar files please place them in the directory *customlib*. This would typical be jdbc drivers, custom threshold classes, etc.

12.2 Developing with Bischeck

Its simple to develop your own service, service item and threshold classes. To develop your own you must follow the interface that exist for each type. For service and service items an abstract class exist with default implementation of most of the methods described in the interfaces.

13 Installation

The latest binary version of Bischeck is available on <http://www.bischeck.org>.

Download the distribution file and follow the steps below to install. Make sure you have root privileges doing this.

```
# tar xzvf bischeck-x.y.z.tgz
# cd bischeck-x.y.z
# chmod 755 install
# ./install -u #Get usage
# ./install #Install default
# service bischeckd start #Redhat/Centos
# /etc/init.d/bischeck start #Debian/Ubuntu
```

To get full list of available options to the install script use -u. By default the *install* script will install Bischeck in directory */opt/socbox/addons/bischeck*, referred to as *\$BISHOME*, and with the ownership of the user id *nagios*. Make sure that the user exist before running install.

The last commands start the Bischeck daemon with the effective user id of the user id set during install, default user *nagios*. The installation will configure bischeckd to start automatically in run level 3, 4 and 5.

The process id of the java process running Bischeck in daemon mode is located in a file, default in */var/tmp/bischeck.pid*. This file is used by the bischeckd script to stop the java program running Bischeck and make sure that only one instance of Bischeck is started on the server.

13.1 Upgrading

From version 0.3.2 upgrading is possible. If you already have an installation download the new version as described in the previous chapter but to upgrade run:

```
# ./install -I /opt/socbox/addons/bischeck -X
```

The upgrade will save the current installation in a directory parallel to the new version named *bischeck_x.y.z*, where x.y.z is the version of the old installation.

The file *migrationpath.txt* describe the supported upgrade paths and what migration scripts that will be run by the *install* script when running in upgrade mode, -X.

If the upgrade is successful Bischeck can be started.

```
# service bischeckd start
```


13.2 Getting started

In the *\$BISHOME/etc* directory there are working examples of all the configuration files. These works as an examples to get started with a simple Bischeck setup of monitoring the response time of a tcp ping on the ssh port, port 22. I think you will find this familiar if you been using Nagios before and the example use the well known check command *check_tcp*. The example assume that you have the Nagios plugins installed. To get the example running you need to do the following steps:

1. Go to the *etc* directory where you installed Bischeck, default is */opt/socbox/addons*.
2. Edit the *bischeck.xml* at line 16 with the path to the *check_tcp* command. If you have *check_tcp* installed in */usr/lib/nagios/plugins* you do not need to edit the file
3. Edit the *servers.xml* file to add correct host name, line 12, and password , line 22, for your NSCA server. Even if you do not have a NSCA server defined Bischeck will still run but with a connection error when trying to send data to NSCA.
4. Now you are done to start Bischeck and checkout what happens in the log file.

```
$ sudo /etc/init.d/bischeckd restart
$ tail -f /usr/tmp/bischeck.log
```

Maybe you ask your self why you should use Bischeck to do a *check_tcp*? Well this is just a simple example, but when you have this running you can check out the *24thresholds.xml* file to get some more understanding about the power of dynamic and adaptive thresholds.

13.3 Integration with pnp4nagios

pnp4nagios can create graph layouts depending on the check command used for the service on the Nagios server. Since Bischeck use a passive check we need to create a unique check command that match the pnp4nagios layout for Bischeck. Create a link in the *libexec* directory on the nagios server:

```
nagios$ ln -s check_dummy check_bischeck
```

When describing the service always use the *check_bischeck* as the check command. The *check_bischeck.php* that control the pnp4nagios layout must be copied to the directory *pnp4nagios/share/templates* on the Nagios server.

13.4 Logging

Bischeck use log4j for log management. The log4j configuration is described in the *log4.properties* file located in the *resources* directory of the Bischeck installation. By default Bischeck writes log information at level INFO to file */var/tmp/bischeck.log*.

14 Command line utilities

There is a number of command line utilities available in Bischeck. All can be ran through the *bischeck* script located in the *\$BISHOME/bin* directory. To use the bischeck script add it to your PATH variable.

```
$ PATH=$BISHOME/bin:$PATH
```

14.1 Run Bischeck

The normal way to run Bischeck is as a daemon using the init.d script *bischeckd*, but is also possible to start Bischeck in continues running mode by executing:

```
$ bischeck Execute -d
```

Running in this way have limitations since the execution will not automatically be placed as a background process and the effective user id will be the user starting the process which may not have all permissions according to the installation. Neither will pid files be updated correctly. For production system always use the init.d script.

```
$ sudo /etc/init.d/bischeckd start
```

or

```
# service bischeckd start
```

For testing purpose it can be good to just run Bischeck once and make sure that every thing is executing as expected. This is done by executing:

```
$ bischeck Execute
```

This will override all scheduling definitions and execute everything directly, but only once.

To show the pid file used for the Bischeck daemon running:

```
$ bischeck ConfigurationManager -p
```

This command is used in the init script bischeckd to retrieve the current pid.

14.2 List Bischeck configuration

Formatted listing of all configuration files is managed by the DocManager class.

To generate the configuration files in html format:

```
$ bischeck DocManager -d <directory name> -t html
```

The files are generated to the directory set with -d. Point your browser to the index.html file in the directory to get links to each configuration file.

To get the same but in plain text:

```
$ bischeck DocManager -d <directory name> -t text
```

If -d is omitted the files file be created in the default directory *bisheckdoc*. The default value for -t is html.

14.3 Validating configuration files

To validate if the xml configuration files are correct the following command will return 0 if correct. Use \$? to see return status.

```
$ bischeck ConfigurationManager -v; echo $?
```

14.4 Twenty4HourThreshold testing

To verify which threshold rule that will used for a specific host, service and service item for a specific date run:

```
$ bischeck threshold.Twenty4HourThreshold -h erphost -s  
orders -i ediorders -d 20111207
```

The above command will show which threshold definition that will be used for host erphost, service orders and serviceitem ediorders on seventh of December 2012. If -d is omitted the current date will be used.

15 Releases

15.1 Release 0.4.3 - 2013-03-27

This release is a minor upgrade.

15.1.1 New feature

- [FR-231] “Add Livestatus as a new server integration alongside nsca and nrdp”. For more information please see 7.1.
- [FR-236] “Extended configuration of threshold for threshold class Twenty4HourThreshold”. For more information please see 8.6.
- The installer script was updated to manage configuration of JMX. The installer add switches for which JMX port to use (-p), which JMX RMI server IP address to bind the port to (-i) and if authentication should be enforced or not (-a). This is related to the [TR-238].
If Bisconf will be used its required to configure JMX RMI support by defining hostname/IP and port.
- The internal surveillance of Bischek that is based on JMX Mbeans was cleaned up. For more information see 11.
- The initial configuration files have changed to an example using a Nagios check command instead of the previous example that used a database connection to mysql. This will be easier to get started with since no JDBC drivers for mysql is required. For more information see the section 13.2.

15.1.2 Bugs fixed and important issues

- [TR-235] “NullPointerException in management of exponential data”
- [TR-237] “Faulty behavior for property lastStatusCacheDumpDir”
- [TR-238] “Bischek will not start without hostname in /etc/hosts”. This is not a bug in Bischek, but more a configuration of JMX to work with remote connections. The install script how been updated to manage the this in a better way.

15.1.3 Upgrading

Release 0.4.3 support upgrade from release 0.4.2.

15.2 Release 0.4.2 - 2012-12-21

Except for the new features introduced and bugs that has been fixed in 0.4.2, there has also been some major work in making Bischeck more stable and performance enhanrunning many host, service and service items configurations. Our goal has been to secure a stable resource utilization of the Java virtual machine (jvm) when running Bischeck over a long period of time.

15.2.1 New feature

- [FR-234] Distribute the start time when running interval based scheduling. This will increase the distribution of service executions especially if Bischeck is configured with many services that is executed on the same interval.
- [FR-233] Support for server integration with Graphite. See section 7.3.
- [FR-232] Support for executing check commands. The check commands that can be used must support an output of performance data. Bischeck will not care about the status that comes from a check command. Instead it will only use the performance data to evaluate its own threshold. This include the new service class *ShellService* and serviceitem class *CheckCommandServiceItem*.
- [FR-224] Support for NRDP through the new server class NRDPServer.
- Related to bug [TR-227] the naming of host, service and serviceitem names has been improved. For more info please see 8.1
- Execution statements and thresholds hour specification where cache data is retrieved as a list, like in a function as `avg(x-y-z[4:10])` and `max(x-y-z[-5M:-15M])`, can now be configured to return a value as long as at lest one index in the range is not null. To support backwards capability the new functionality will only be used if the property *notFullListParse* is set to true in the properties.xml. The default value is false. Read more about this option 5.3.
- There has been some discussion about what Nagios state should be sent if the a the returned execute statement of a service item is null. In previous releases this has been hard coded to OK, but now its possible to define it by setting the property *stateOnNull*. The property can be set to an integer 0,1,2 or 3 or to a string OK, WARNING, CRITICAL or UNKNOWN. The default is UNKNOWN.
- When a service class get an exception when creating a connection the previous versions did not save any data to the cache. If the property *saveNullOnConnectionError* you will now get a null value inserted into the cache when a connection exception is thrown. For backwards compatibility the default value of the property is false.

- More mathematical functions like `multNull` and `divNull` that support null value as part of the calculation and can be used with functions that take a list of numbers to manage calculation where cache data may be null.

15.2.2 Bugs fixed and important issues

- [TR-227] “Cache parser do not work for host, service or serviceitems if the name include 0 (zero)” has been resolved.
- [TR-228] “Threshold factory return wrong threshold definition if service and serviceitem name is the same for different hosts” has been resolved.
- [TR-229] “When using service `ShellService` the number of open files limit will be reached” has been resolved.
- [TR-230] “NRDP submissions all come in as OK” has been resolved.
- Fixed migration script from 0.4.0 to copy *etc* directory content correctly. Changes in the file *urlservices.xml* will be overwritten. Existing 0.4.0 configuration will still be available in the previous version backup directory, `bischeck_0.4.0`.

15.2.3 Upgrading

Release 0.3.3, 0.4.0 and 0.4.1 is supported for upgrade to 0.4.2. The upgrading is NOT applicable for release candidate.

15.3 Release 0.4.1 - 2012-10-01

15.3.1 New feature

- [FR-224] **Beta** support for NRDP through the new server class `NRDPService`.
- **Beta** support for executing check commands. The check commands that can be used must support an output of performance data. Bischeck will not care about the status that comes from a check command. Instead it will only use the performance data to evaluate its own threshold. This include the new service class *ShellService* and serviceitem class *CheckCommandServiceItem*.
- [FR-223] Instrumentation metrics provided by <http://metrics.codahale.com/> has been implemented to enable fine grain real time process measuring through JMX. This is implemented for
 - All external interfaces to measure response time
 - Service execution time
 - Threshold processing time, etc.

15.3.2 Bugs fixed and important issues

- [TR-225] “A service url that has no equivalent in urlservices.xml generate a NullPointerException” has been resolved.
- [TR-226] “For thresholds with more the two decimals will not be correctly validated” has been fixed. The bug was related to when Bischeck was used with data points with many decimals, like the result from a ping. Bischeck would strip of all decimals except 2 with caused small values to become 0. Now Bischeck will determine the number of decimals used by the collected data point and use it when formatting the calculated threshold, warning and critical value. This will hopefully make data presented in Nagios UIs and graphs like pnp4nagios look better.

15.3.3 Upgrading

Release 0.4.0 is supported for upgrade to 0.4.1.

15.4 Release 0.4.0 - 2012-08-31

15.4.1 New feature

- [FR-197] Support for different and multiple integration with different surveillance and monitoring systems. With version 0.4.0 Bischeck is not limited to send data to Nagios. It can now send the data to multiple Nagios servers and to other servers like OpenTSB. This is done by moving server formatting and protocol to server integration classes that implements the interface `com.ingby.socbox.bischeck.servers.Server`. The server integration is described in the xml configuration file `servers.xml`. This also means that that some Nagios NSCA specific properties previous configured in `properties.xml` has been moved to the `servers.xml` file in the NSCA section. The OpenTSDB server class should be regarded as beta.
- [FR-204] The bischeck cache will be saved when the bischeck daemon is shutdown and reloaded on bischeck startup. Keeping the cache persistent between restarts is important since 0.4.0 support time based cache retrieval. The limitations is currently that if the Bischeck daemon is killed by a signal that can not be caught or the daemon crash the data will not be saved. This will be improved in future versions.
- [FR-202] The implementation of running bischeck once, in a none daemon mode, is changed so the same code is used as running in daemon mode. The only difference is that the initialization of triggers are different so all service items are just ran directly and and just once.
- [FR-218] The bischeck daemon can now reload the configuration without a process restart. This is support through the JMX operation “reload”. The feature will limit the need of operating system access and authorization.

- [FR-219] Bischeck can now retrieve state and performance data from a Nagios server supporting livestatus. With the service class `LivestatusService` a connection is set up over livestatus and with the and serviceitem class `LivestatusServiceItem` state and/or performance data can be retrieved from the a Nagios service. This can be useful when when creating virtual services in Bischeck or used in complex thresholds.
- [FR-220] Bischeck now support one additional scheduling method where scheduling can be defined to run a service after a different service has executed. This can be useful when a service is depending on data for another service for its thresholds or execution statement.
- [FR-221] Cache retrieval is now support by using a time offset to find the nearest cache element to the time offset.
- Cache data can be retrieved as a list of elements based both on index and time.
- Support for additional mathematical functions like average, min and max calculations on list of elements.
- Bischeck can now support the usage of cached data in an execution statement of a serviceitem. This is typical useful when a serviceitem execute statement is depending on other service data. For example in a SQL query string:

```
select value from table1 where id = host1-web-state[0]
and createdate = '%%yyyy-MM-dd%'");
```

- Added support for other Linux distributions then Redhat based. Bischeck should now install on Debian 6 and Ubuntu 10/11.
- Configuration listing. The configuration listing has been moved from the ConfigurationManager class to the DocManager class. Currently html and text listing is supported. The generated configuration data will by default placed in the *bischeck-doc* directory.
- A configured service can be configured not to send its data to a the configured monitoring servers like Nagios. This can be useful if the service is just to be used to create virtual services or just to be used as thresholds.
- The bischeck script now support JMX authentication. The authentication files are located in the etc directory and named `jmxremote.password` and `jmxremote.access`. Default is to that authentication is disabled by the system property `-Dcom.sun.management.jmxremote.authenticate=false`. To enable authentication set the property to true. For more info about JMX see <http://www.oracle.com/technetwork/java/javase/tech/javamanagement-140525.html>.

15.4.2 Bugs fixed and important issues

- The Twenty4Thresholds class was in previous version not doing a correct linear equation calculation if a expression based threshold was defined. Lets illustrate the errors with this example from the 24thresholds.xml configuration file having a mix with static and expression based thresholds.

```

. . . .
<!-- 12:00 -->
<hour>7000</hour>
<!-- 13:00 -->
<hour>testhost-testservice-testitem[1] / 3</hour>
<!-- 14:00 -->
<hour>testhost-testservice-testitem[1] / 2</hour>
<!-- 15:00 -->
<hour>testhost-testservice-testitem[1] + 1000 </hour>
<!-- 16:00 -->
<hour>12000</hour>
. . . .

```

In the previous version the threshold value between 12:00 and 13:00 would be null since it was a mix of static and expression based thresholds. And between 15:00 and 16:00 the threshold would have been calculated as “testhost-testservice-testitem[1] + 1000” independent of the time between 15:00 and 16:00.

Now the linear equation will correctly be calculated with any mix of static and expression based definitions. In the above example the calculated threshold for 12:20 will now be:

$$20 * ((\text{testhost-testservice-testitem}[1] / 3) - 7000) / 60 + 7000$$

This fix will improve the correctness and also the capability of threshold adaptivity.

- The Service interface has a number of new methods that should been there from the beginning. If you developed any service class you need to add these, but if you just inherited ServiceAbstract its fixed for you. The new methods are:

```

public NAGIOSSTAT getLevel();
public void setLevel(NAGIOSSTAT level);
public boolean isConnectionEstablished();
public void setConnectionEstablished(boolean connected);
public Boolean isSendServiceData();
public setSendServiceData(Boolean sendServiceData);

```

- Property *cacheClear* is renamed to *thresholdCacheClear*.

- All the nsca related properties has been moved from properties.xml to servers.xml when used for the NSCAServer class. The new property names has also gone through some minor changes. When upgrading a manual update is needed of the servers.xml file with the current setting of nsca related properties in properties.xml. Recommended that these are later removed.
- All JAXB generated configuration classes now support serialization.
- Quartz jar is upgraded from 2.0.1 to 2.1.5.
- [TR-216] “Shutdown is automatic triggered”
- [TR-217] “Configuration Manager initialization failed with java.lang.NullPointerException”
- [TR-207] “sudo in bischeckd script cause problem at boot”

15.4.3 Upgrading

Release 0.3.3 and 0.4.0_RC2 are supported for upgrade to 0.4.0.

15.5 Release 0.3.3 - 2011-11-14

15.5.1 New features

- Bischeck are no longer limited to just be integrated with a single Nagios server over the NSCA protocol. Now is it possible to integrate with multiple monitoring servers over different protocols. Currently Nagios/NSCA and OpenTSDB is support. To enable this a new class component called Server has been introduced. The class is responsible for communication and formatting for the specific monitoring server it integrate against. A new configuration file, server.xml is used for configuration of server integration.

15.5.2 Bugs fixed and important issues

- [TR-207] “sudo in bischeckd script cause problem at boot”
- [TR-214] “Threshold object cache is no cleared”

15.6 Release 0.3.2 - 2011-07-29

15.6.1 New features

- The configuration system has been completely rewritten and now us xml based configuration files. Each configuration file has a corresponding xsd file that can be used for verifications. The dependencies to sqlite3 has been deprecated and is just part of this release to support upgrade.

- The scheduling of services and its related serviceitem(s) has been rewritten to support different scheduling policies per service instead of earlier versions of fixed interval scheduling. With 0.3.2 each service can have one to many schedule tags in bischeck.xml configuration file. For more info please see the chapter 9 on page 35.

15.6.2 Bugs fixed and important issues

- The active attribute on Hosts, Services and Serviceitem has been removed.
- The interface com.ingby.socbox.bischeck.threshold.Threshold has a new signature on the method init(). This method now throws Exception.
- The Service interface has two additional methods, setSchedules() and getSchedules().
- The Service interface has changed the signature of getServicesItems() to return Map instead of HashMap.
- buildr has been replaced by ant as the build management system.

15.7 Release 0.3.1 - 2011-04-08

15.7.1 New features

- The ServiceFactory class now use a property table, urlservice, to map what Service class should be instantiate for a specific url schema. The url schema is the key. The current default mapping are:
 - jdbc -> JDBCService
 - bischeck -> LastCacheService
- The ServiceItemFactory class use an additional field, serviceitemclass, in the items configuration table to determine what ServiceItem class to instantiate.
- Calendar in Bischeck follows the ISO 8601 date standard by default. This means that the first day in the week is Monday, day 2 according to java.util.Calendar, and that the first week of the year must have a minimum of 4 days. The importance of this is to get the week numbering correct that is used in the configuration in TwentyFourHourThreshold class, but day one (1) in the week is still Sunday when defining the tag dayofweek in 24thresholds.xml. The setting can be overridden by setting the properties “*mindaysinfirstweek*” (default 4) and “*firstdayofweek*” (default 2) in the properties.xml file.
- If no threshold class has been specified, null in the thresholdclass field in the items table, Bischeck will instantiate the “empty” class DummyThreshold.

- For all class configuration of Service, ServiceItem and Threshold its now possible to specify the class name without the package path if the class is part of the Bischeck distribution.
- Clean up of the exception handling process when starting Bischeck. Now the execution should not start if there are configuration issues with missing classes for Service, ServiceItem and Threshold.

15.7.2 Bugs fixed and important issues

- N/A

15.7.3 Upgrade issues

- Upgrade by doing a fresh installation, but first save the old installation directory. After saving the old installation do a new install. Then copy the files bischeck.conf and 24threshold.conf from old to new installation dir.
- The field serviceitemclass (varchar(256)) in table items in configuration database bischeck.conf must be manual added and populated with the right Service class name. If corresponding service is jdbc:// set the field serviceitemclass to SQLServiceItem and if the service is bischeck:// set the field to CalculateOnCache.
- To add the column:

```
$ sqlite3 bischeck.conf sqlite> ALTER TABLE items ADD  
    COLUMN serviceitemclass varchar(256);
```

- Update the serviceitemclass for all rows in items:
 - sqlite> update items set serviceitemclass='SQLServiceItem' where
 - sqlite> update items set serviceitemclass='CalculateOnCache' where
- Add the new table url2service in database bischeck.conf.

```
$ cat << EOF | sqlite3 bischeck.conf  
drop table IF EXISTS urlservice;  
create table urlservice(key varchar(128), value varchar(256));  
insert into urlservice values ("jdbc","JDBCService");  
insert into urlservice values ("bischeck","LastCacheService");  
EOF
```

- Copy all file located in the old installation customlib directory to the customlib directory in the new installation.

15.8 Release 0.3.0 - 2011-03-03

15.8.1 New features

- This is the first binary distribution, but should be regarded as a beta version.

15.8.2 Bugs fixed and important issues

- N/A

16 System requirements

Bischeck should run on any operating system that supports Java 6. The installation script and init scripts are supported on Redhat and Debian equivalent Linux distributions. Running on none Linux operating system has not been tested.

The following jar packages are distributed as part of the Bischeck distribution. All these packages have their own open source licenses.

- commons-lang-2.5.jar - <http://commons.apache.org/lang/>
 - <http://www.apache.org/licenses/LICENSE-2.0.html>
- commons-cli-1.2.jar - <http://commons.apache.org/cli/>
 - <http://www.apache.org/licenses/LICENSE-2.0.html>
- log4j-1.2.16.jar - <http://logging.apache.org/log4j/Apacheproject>
 - <http://www.apache.org/licenses/LICENSE-2.0.html>
- quartz-2.1.5.jar - <http://www.quartz-scheduler.org/>
 - <http://www.apache.org/licenses/LICENSE-2.0.html>
- jsendnsca-2.0.1.jar - <http://code.google.com/p/jsendnsca/>
 - <http://www.apache.org/licenses/LICENSE-2.0.html>
- slf4j-api-1.6.0.jar - <http://www.slf4j.org/>
 - <http://www.opensource.org/licenses/mit-license.php>
- slf4j-log4j12-1.6.0.jar - <http://www.slf4j.org/>
 - <http://www.opensource.org/licenses/mit-license.php>
- jep-2.3.1.jar - <http://sourceforge.net/projects/jep/>
 - <http://www.gnu.org/licenses/gpl-2.0.html>
- json-lib-2.4 - <http://json-lib.sourceforge.net/>
 - <http://www.apache.org/licenses/LICENSE-2.0.html>
- ezmorph-1.0.6.jar - <http://ezmorph.sourceforge.net/>
 - <http://www.apache.org/licenses/LICENSE-2.0.html>

- commons-beanutils-1.8.3.jar - <http://commons.apache.org/beanutils/>
 - <http://www.apache.org/licenses/LICENSE-2.0.html>
- commons-collections-3.2.1.jar - <http://commons.apache.org/collections/>
 - <http://www.apache.org/licenses/LICENSE-2.0.html>
- metrics-core.jar - <https://github.com/codahale/metrics>
 - <http://www.apache.org/licenses/LICENSE-2.0.html>

All jar files distributed as part of Bischeck are located in the lib directory.

17 Bischeck license

Bischeck is licensed under GNU license version 2. For more info please visit <http://www.gnu.org/licenses/gpl-2.0.html>

18 Bug reports and feature requests

Please submit bug reports and feature requests on www.bischeck.org in the Forge section.

19 Credits

Thanks to all people who has developed all the great software that Bischeck depends on, and especially all who made Nagios and the Nagios community a success. A special thanks goes to Peter Johansson and his colleagues at DHL Freight Sweden for sponsoring the development and providing ideas and important feedback.